

Writing a Mini-Lisp Compiler

Carlo Hamalainen

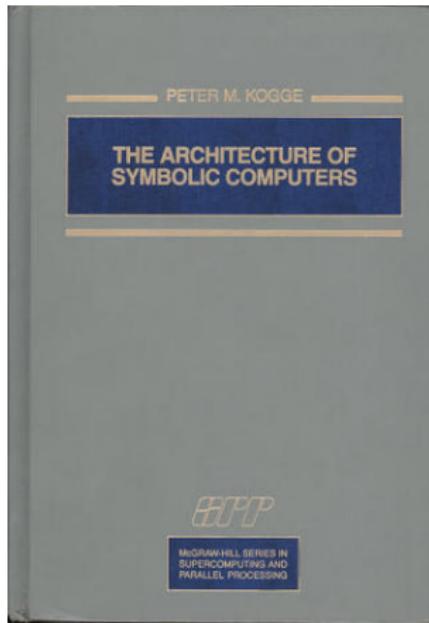
carlo-hamalainen.net

28 May 2013

About 13 years ago I did an BSc/BInfTech combined degree.

A few of the IT subjects were fun:

- Computational complexity. NP-complete, etc.
- Operating systems. We wrote our own memory pager in C.
- Functional programming. Gofer on MS-DOS.
- Compilers. We wrote a compiler for a mini-Java language, in C.



The Architecture of Symbolic Computers (Peter M. Kogge, 1991)

The Architecture of Symbolic Computers (Peter M. Kogge, 1991)

<http://blog.fogus.me/2012/07/25/some-lisp-books-and-then-some/>

“The best Lisp book I’ve ever read and could ever possibly read.”

<http://www.loper-os.org/?p=13>

“The Architecture of Symbolic Computers (Peter M. Kogge) is quite possibly the most useful resource I have come across in my quest thus far. If you are interested in Lisp Machine revival, non-von Neumann computation, or the dark arts of the low-level implementation of functional programming systems, you will not be disappointed.”

Here's the plan:

- ① Define a virtual machine: memory, registers, opcodes.
- ② Write an emulator.
- ③ Define a mini-Lisp language (let, lambdas, etc).
- ④ Write a compiler for mini-Lisp to the virtual machine.

IBM 704/709 - circa 1950

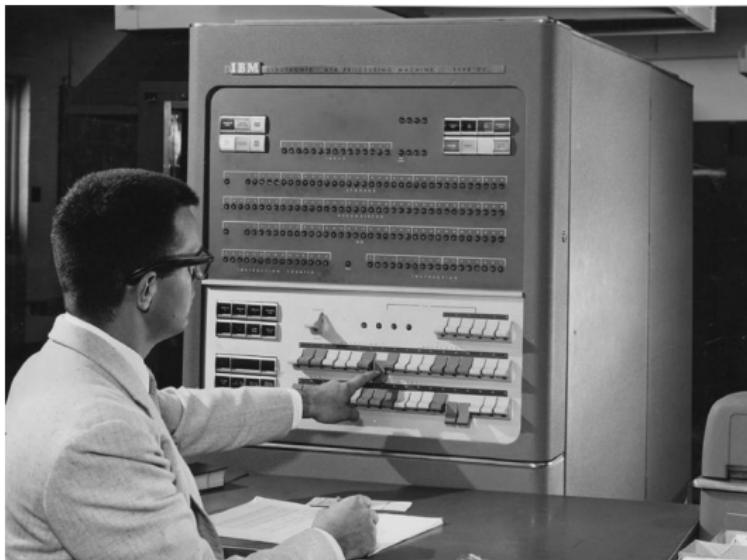
16K per box!



<http://www.computer-history.info/Page4.dir/pages/IBM.704.dir/>

IBM 704/709 - circa 1950

Operator console.



<http://www.computer-history.info/Page4.dir/pages/IBM.704.dir/>

IBM 701 console



The IBM 704 had 36-bit machine words:

- address part: 15 bits
- decrement part: 15 bits
- prefix part: 3 bits
- tag part: 3 bits

Corresponding instructions:

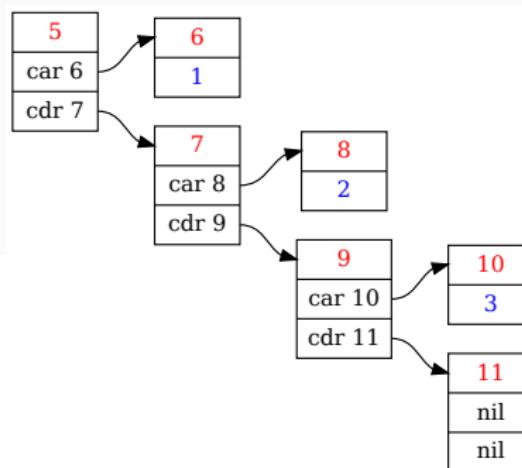
- CAR: Contents of Address part of Register
- CDR: Contents of Decrement part of Register

Think of CAR as head and CDR as tail. Store linked lists with ease.

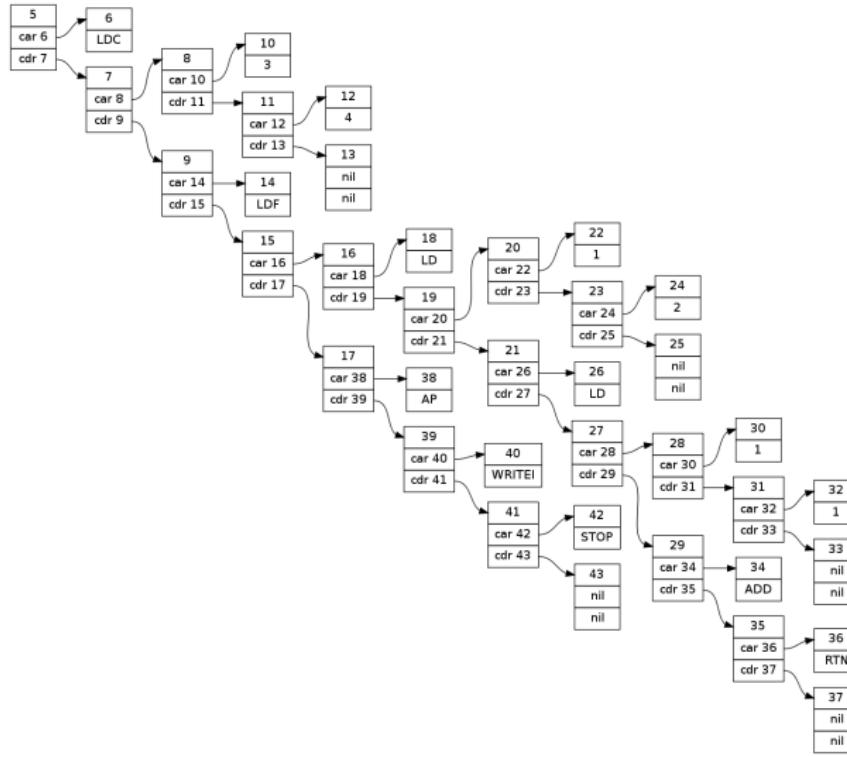
Store this list:

[1 , 2 , 3]

address	value
<hr/>	
5	(6 , 7)
6	1
7	(8 , 9)
8	2
9	(10 , 11)
10	3
11	(nil , nil)



```
[LDC , [3 , 4] , LDF , [LD , [1 , 2] , LD , [1 , 1] , ADD , RTN] ,  
AP , WRITEI , STOP ,]
```



A linked list can represent a stack:

- push: insert at front, return new head.
- pop: remove from front, return cdr (tail) of original head.

We can run simple computations if we have two things:

- A stack for temporary results (like an accumulator register); and
- a pointer to the next instruction to execute.

- S: temporary results stack.
- C: program counter for currently executing code.
- LDC: LoaD Constant
- ADD: integer ADDition
- STOP: STOP the machine

Add 3 and 4:

1. S: [] C: [LDC , 3, LDC , 4, ADD , STOP]
2. S: [3] C: [LDC , 4, ADD , STOP]
3. S: [4, 3] C: [ADD , STOP]
4. S: [7] C: [STOP]

```
LDC    = 100
ADD    = 101
STOP   = 102

S = []
C = [LDC, 3, LDC, 4, ADD, STOP]

def opcode_LDC(S, C):
    assert C[0] == LDC
    C.pop(0)                      # pop LDC
    S.insert(0, C.pop(0))          # push value onto S

def opcode_ADD(S, C):
    assert C[0] == ADD
    C.pop(0)                      # pop ADD
    val1 = S.pop(0)                # first value
    val2 = S.pop(0)                # second value
    S.insert(0, val1 + val2)       # push result
```

Here's the machine emulator:

```
S = []
C = [LDC, 3, LDC, 4, ADD, STOP]

while True:
    if C[0] == STOP: break
    elif C[0] == LDC:
        opcode_LCD(S, C)
    elif C[0] == ADD:
        opcode_ADD(S, C)
    else:
        assert False
```

Summary of the machine so far:

- Memory: a cell contains an integer or (car, cdr) tuple.
- Registers: S (stack), C (code).
- Opcodes: ADD, LDC, STOP.

This doesn't do much — we need functions of some sort.

Lambda expressions:

- Python: `(lambda x, y: x + y)(3, 4)`
- Haskell: `(\x y -> x + y) 3 4`
- Haskell': `(\x y -> (+) x y) 3 4`
- Lisp: `((LAMBDA (x y) (+ x y)) (3 4))`

Problems:

- Can't store "x", so remember `x` = first parameter, etc.
- Save environment before call: new register D (dump).
- Store arguments: new register E (environment).
- New opcodes:
 - LD: LoaD identifier. e.g. [LD, 1] = load value of `x`.
 - LDF: LoaD Function. Push a closure onto S.
 - AP: APply function. Evaluate a closure.
 - RTN: ReTuRN from function. Restore state.

```
S: []
E: []
C: [ LDC, [3, 4] , LDF, [LD, 2, LD, 1, ADD, RTN] , AP , STOP]
D: []
```

```
S: [ [3, 4] ]
E: []
C: [ LDF, [LD, 2, LD, 1, ADD, RTN] , AP , STOP]
D: []
```

```
S: [ [LD, 2, LD, 1, ADD, RTN] , [3, 4] ] # a closure
E: []
C: [ AP , STOP]
D: []
```

```
S: [ [LD, 2, LD, 1, ADD, RTN] , [3, 4] ] # a closure
E: []
C: [ AP , STOP]
D: []
```

```
S: []
E: [ [3, 4] ]
C: [LD, 2, LD, 1, ADD, RTN]
D: [[S=[] , E=[] , C=[STOP]]]
```

```
S: [4]
E: [ [3, 4] ]
C: [LD, 1, ADD, RTN]
D: [[S=[] , E=[] , C=[STOP]]]
```

S: [3 , 4]

E: [[3, 4]]

C: [ADD, RTN]

D: [[S= [] , E= [] , C=[STOP]]]

S: [7]

E: [[3, 4]]

C: [RTN]

D: [[S= [] , E= [] , C=[STOP]]]

S: [7]

E: []

C: [STOP]

D: []

SECD machine summary so far:

- Memory: a cell contains an integer or (car, cdr) tuple.
- Registers: S (stack), E (environment), C (code), D (dump).
- Opcodes: ADD, LDC, STOP LD, LDF, RTN, AP.
- Can evaluate lambda expressions.

SECD opcodes

```
ADD      # integer addition
MUL      # integer multiplication
SUB      # integer subtraction
DIV      # integer division

NIL      # push nil pointer onto the stack
CONS     # cons the top of the stack onto the next list
LDC      # push a constant argument
LDF      # load function
AP       # function application
LD       # load a variable
CAR      # value of car cell
CDR      # value of cdr cell

DUM      # setup recursive closure list
RAP      # recursive apply
```

SECD opcodes contd.

```
JOIN    # C = pop dump
RTN     # return from function
SEL     # logical selection (if/then/else)
NULL    # test if list is empty

WRITEI  # write an integer to the terminal
WRITEC # write a character to the terminal

READC  # read a single character from the terminal
READI  # read an integer from the terminal

ZEROP   # test if top of stack = 0 [nonstandard opcode]
GTOP    # test if top of stack > 0 [nonstandard opcode]
LTOP    # test if top of stack < 0 [nonstandard opcode]

STOP    # halt the machine
```

Final steps...

- Define a simple Lisp language.
- Write a function `compile(x)` to translate Lisp forms into SECD machine code.

Square-bracket Lisp

```
# 3 + 4
[ADD, 3, 4]

# lambda x, y: x + y
[LAMBDA, ['x', 'y'], [ADD, 'x', 'y']]

# (lambda x, y: x + y)(3, 4)
[[LAMBDA, ['x', 'y'], [ADD, 'x', 'y']], 3, 4]

# x = 5; y = 7; x - y
[LET, ['x', 'y'], [5, 7], [SUB, 'x', 'y']]
```

Square-bracket Lisp

```
[ADD, 3, 4]

>>> compile([ADD, 3, 4])
[LDC, 4, LDC, 3, ADD, STOP]

def compile(x):
    if x[0] == ADD:
        x.pop(0)
        arg1 = x.pop(0)
        arg2 = x.pop(0)
        return [LDC, compile(arg1), \
                LDC, compile(arg2), \
                ADD] \
               + compile(x)
    elif ...
```

Square-bracket Lisp

```
# lambda x, y: x + y
[LAMBDA, [ 'x', 'y' ], [ADD, 'x', 'y']]]

>>> compile([LAMBDA, [ 'x', 'y' ], [ADD, 'x', 'y']])
[LDF, [LD, [1, 2], LD, [1, 1], ADD, RTN]]
```

Square-bracket Lisp

```
# (lambda x, y: x + y)(3, 4)
[[LAMBDA, ['x', 'y'], [ADD, 'x', 'y]], 3, 4]

>>> c = [[LAMBDA, ['x', 'y'], [ADD, 'x', 'y]], 3, 4]
>>> compile(c)
[NIL, LDC, 4, CONS, LDC, 3, CONS,
 LDF, [LD, [1, 2], LD, [1, 1], ADD, RTN],
 AP,
 STOP]
```

Square-bracket Lisp

```
# x = 5; y = 7; x - y
[LET, [ 'x' , 'y' ] , [5, 7] , [SUB , 'x' , 'y' ]]

>>> c = [LET, [ 'x' , 'y' ] , [5, 7] , [SUB , 'x' , 'y' ]]
>>> compile(c)
[NIL, LDC, 7, CONS, LDC, 5, CONS,
 LDF, [LD, [1, 2], LD, [1, 1], SUB, RTN],
 AP,
 STOP]
```

Length of a list?

```
> let f x m = if null x
              then m
              else f (tail x) (m + 1)
> f [1, 2, 3] 0
3
```

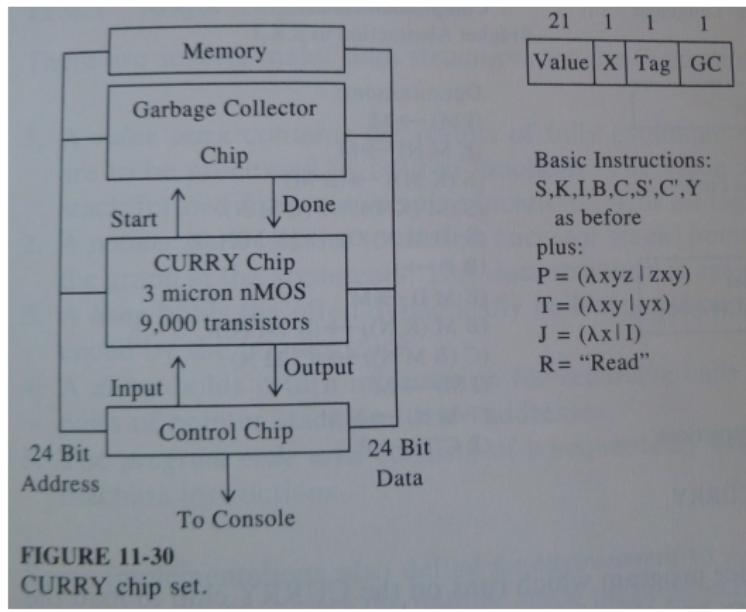
```
[LETREC, [ 'f' ], [[LAMBDA, [ 'x' , 'm' ] ,
  [IF, [NULL, 'x'],
   'm',
   [ 'f' , [CDR, 'x'] , [ADD, 'm', 1]]]]],
 [ 'f' , [LIST, 1, 2, 3], 0]]
```

Square-bracket Lisp:

- Enough to write basic programs: LET, LAMBDA, IF.
- Recursion using LETREC.
- No destructive assignment: pure Lisp.
- Closures can be evaluated in the future, always same value.
- Code as data.
- Lambda expressions \leftrightarrow S-expressions.
- Relies on machine to do garbage collection (say what?).

The CURRY chip (Ramsdell, 1986)

- 0.4Mhz
- CPU chip: about 9000 transistors, performed graph reduction (combinators S, K, I).
- Memory chip: garbage collection.
- Timing chip: overall timing and control console.



MultiLISP: parallel Lisp

- Pure lambda calculus \Rightarrow opportunities for parallelism.
- Concert machine at MIT: 24-way Motorola 68000-based shared-memory multiprocessor (early 80s?).
- MultiLISP \rightarrow SECD-like abstract machine called MCODE.
- MCODE interpreter (3000 lines of C).
- Unfortunately slow due to MCODE, but parallel benefits could be measured.
- Used *futures* for evaluation; see also Swift:
<http://www.ci.uchicago.edu/swift/main/>

Symbolics Lisp Machine



Symbolics Lisp Machine

- Instruction set: stack machine, similar to SECD.
- Hardware support:
 - virtual memory
 - garbage collection
 - data type testing

Symbolics Lisp Machine keyboard



modifier keys: HYPER SUPER META CONTROL

My SECD emulator and mini-Lisp compiler:
<https://github.com/carlohamalainen/pysecd>